

**This document will guide you regarding how to install multinode cloudera hadoop cluster cdh5.4.0 without Cloudera manager.**

**In this tutorial I have used 2 Centos 6.6 virtual machines viz. master.hadoop.com & slave.hadoop.com.**

**Prerequisites:**

**CentOS 6.X**

**Open JDK 1.7:-**

```
sudo yum install java-1.7.0-openjdk-devel
```

**Master machine** – master.hadoop.com (ip of first machine)

Daemons that we are going to install on master are :

Namenode

HistoryServer

**Slave machine** – slave.hadoop.com (ip of second machine)

Daemons that we are going to install on master are :

Resource Manager (Yarn)

Node-manager

Secondary Namenode

Datanode

**Important configuration before proceeding further: please add both the hostname and ip information to `/etc/hosts` file on each host.**

```
[root@master ~]# cat /etc/hosts
127.0.0.1 localhost localhost.localdomain localhost4 localhost4.localdomain4
::1      localhost localhost.localdomain localhost6 localhost6.localdomain6
192.168.111.130    master.hadoop.com
192.168.111.131    slave.hadoop.com
```

```
[root@slave ~]# cat /etc/hosts
127.0.0.1 localhost localhost.localdomain localhost4 localhost4.localdomain4
::1      localhost localhost.localdomain localhost6 localhost6.localdomain6
192.168.111.130    master.hadoop.com
192.168.111.131    slave.hadoop.com
```



Please verify that both the hosts are ping'able from each other

Also,

Please stop the firewall and disable the selinux.

**To stop firewall in centos :**

```
service iptables stop && chkconfig iptables off
```

**To disable selinux :**

```
vim /etc/selinux/config
```

once file is opened, please verify that "SELINUX=disabled" is set.

### **1.Date should be in sync**

Please make sure that master and slave machine's date is in sync, if not please do it so by configuring NTP.

### **2.Passwordless ssh must be setup from master → slave**

**To setup passwordless ssh follow below procedure :**

#### **2a. Generate rsa key pair using ssh-keygen command**

```
[root@master conf]# ssh-keygen  
Generating public/private rsa key pair.
```

```
Enter file in which to save the key (/root/.ssh/id_rsa):
```

```
/root/.ssh/id_rsa already exists.
```

```
Overwrite (y/n)?
```

## 2b. Copy generated public key to slave.hadoop.com

```
[root@master conf]# ssh-copy-id -i ~/.ssh/id_rsa.pub root@slave.hadoop.com
```

Now try logging into the machine, with “ssh ‘root@slave.hadoop.com’”, and check in:

```
.ssh/authorized_keys
```

to make sure we haven’t added extra keys that you weren’t expecting.

## 2c. Now try connecting to slave.hadoop.com using ssh

```
[root@master conf]# ssh root@slave.hadoop.com
```

```
Last login: Fri Apr 24 14:20:43 2015 from master.hadoop.com
```

```
[root@slave ~]# logout
```

```
Connection to slave.hadoop.com closed.
```

```
[root@master conf]#
```

That’s it! You have successfully configured passwordless ssh between master and slave node.

## 3. Internet connection

Please make sure that you have working internet connection, as we are going to download CDH packages in next steps.

#### **4. Install cdh repo**

##### **4a. download cdh repo rpm**

```
[root@master ~]# wget http://archive.cloudera.com/cdh5/one-click-install/redhat/6/x86_64/cloudera-cdh-5-0.x86_64.rpm
```

##### **4b. install cdh repo downloaded in above step**

```
[root@master ~]# yum --nogpgcheck localinstall cloudera-cdh-5-0.x86_64.rpm
Loaded plugins: fastestmirror, refresh-packagekit, security
setting up Local Package Process
....
Complete!
```

##### **4c. do the same steps on slave node**

```
[root@slave ~]# wget http://archive.cloudera.com/cdh5/one-click-install/redhat/6/x86_64/cloudera-cdh-5-0.x86_64.rpm
```

```
[root@slave ~]# yum --nogpgcheck localinstall cloudera-cdh-5-0.x86_64.rpm
```

Loaded plugins: fastestmirror, refresh-packagekit, security

Setting up Local Package Process

.....

Complete!

## 5. Install and deploy ZooKeeper.

```
[root@master ~]# yum -y install zookeeper-server
```

Loaded plugins: fastestmirror, refresh-packagekit, security

Setting up Install Process

.....

Complete!

### 5a. create zookeeper dir and apply permissions

```
[root@master ~]# mkdir -p /var/lib/zookeeper
```

```
[root@master ~]# chown -R zookeeper /var/lib/zookeeper/
```

### 5b. Init zookeeper and start the service

```
[root@master ~]# service zookeeper-server init
```

No myid provided, be sure to specify it in /var/lib/zookeeper/myid if using non-standalone

```
[root@master ~]# service zookeeper-server start
```

JMX enabled by default

Using config: /etc/zookeeper/conf/zoo.cfg

Starting zookeeper ... STARTED

### 6. Install namenode on master machine

```
yum -y install hadoop-hdfs-namenode
```

### 7. Install secondary namenode on slave machine

```
yum -y install hadoop-hdfs-secondarynamenode
```

### 8. Install resource manager on slave machine

```
yum -y install hadoop-yarn-resourcemanager
```

## 9. Install nodemanager, datanode & mapreduce on slave node

```
yum -y install hadoop-yarn-nodemanager hadoop-hdfs-datanode hadoop-mapreduce
```

## 10. Install history server and yarn proxyserver on master machine

```
yum -y install hadoop-mapreduce-historyserver hadoop-yarn-proxyserver
```

## 11. On both the machine you can install hadoop-client package

```
yum -y install hadoop-client
```

**Now we are done with the installation, it's time to deploy HDFS!**

### 1. On each node, execute below commands :

```
[root@master ~]# cp -r /etc/hadoop/conf.empty /etc/hadoop/conf.my_cluster  
[root@master ~]# alternatives --install /etc/hadoop/conf hadoop-conf /etc/hadoop/conf.my_cluster  
50  
[root@master ~]# alternatives --set hadoop-conf /etc/hadoop/conf.my_cluster
```

```
[root@slave ~]# cp -r /etc/hadoop/conf.empty /etc/hadoop/conf.my_cluster  
[root@slave ~]# alternatives --install /etc/hadoop/conf hadoop-conf /etc/hadoop/conf.my_cluster 5  
0
```



```
[root@slave ~]# alternatives --set hadoop-conf /etc/hadoop/conf.my_cluster
```

## 2. Let's configure hdfs properties now :

Goto /etc/hadoop/conf/ dir on master node and edit below property files:

### 2a. vim /etc/hadoop/conf/core-site.xml

Add below lines in it under <configuration> tag

```
<property>
<name>fs.defaultFS</name>
<value>hdfs://master.hadoop.com:8020</value>
</property>
```

### 2b. vim /etc/hadoop/conf/hdfs-site.xml

```
<property>
<name>dfs.permissions.superusergroup</name>
<value>hadoop</value>
</property>
<property>
<name>dfs.namenode.name.dir</name>
<value>file:///data/1/dfs/nn,file:///nfsmount/dfs/nn</value>
```

```
</property>
<property>
<name>dfs.datanode.data.dir</name>
<value>file:///data/1/dfs/dn,file:///data/2/dfs/dn,file:///data/3/dfs/dn,file:///data/4/dfs/dn</value>
</property>
<property>
<name>dfs.namenode.http-address</name>
<value>192.168.111.130:50070</value>
<description>
The address and the base port on which the dfs NameNode Web UI will listen.
</description>
</property>
```

### 3. scp core-site.xml and hdfs-site.xml to slave.hadoop.com at /etc/hadoop/conf/

```
[root@master conf]# scp core-site.xml hdfs-site.xml slave.hadoop.com:/etc/hadoop/conf/
core-site.xml                                100% 1
001   1.0KB/s  00:00
hdfs-site.xml                                100% 1
669   1.6KB/s  00:00
[root@master conf]#
```

### 4. Create local directories:

On master host run below commands:

```
mkdir -p /data/1/dfs/nn /nfsmount/dfs/nn  
chown -R hdfs:hdfs /data/1/dfs/nn /nfsmount/dfs/nn  
chmod 700 /data/1/dfs/nn /nfsmount/dfs/nn  
chmod go-rx /data/1/dfs/nn /nfsmount/dfs/nn
```

**On slave host run below commands:**

```
mkdir -p /data/1/dfs/dn /data/2/dfs/dn /data/3/dfs/dn /data/4/dfs/dn  
chown -R hdfs:hdfs /data/1/dfs/dn /data/2/dfs/dn /data/3/dfs/dn /data/4/dfs/dn
```

**5. Format the namenode :**

```
sudo -u hdfs hdfs namenode -format
```

**6. Start hdfs services**

**Run below commands on master and slave**

```
for x in `cd /etc/init.d ; ls hadoop-hdfs-*` ; do service $x start ; done
```

**7. Create hdfs tmp dir**

**Run on any of the hadoop node**

```
[root@slave ~]# sudo -u hdfs hadoop fs -mkdir /tmp  
[root@slave ~]# sudo -u hdfs hadoop fs -chmod -R 1777 /tmp
```

**Congratulations! You have deployed hdfs successfully**



## **Deploy Yarn**

### **1. Prepare yarn configuration properties**

replace your `/etc/hadoop/conf/mapred-site.xml` with below contents on master host

```
[root@master conf]# cat mapred-site.xml
```

```
<configuration>  
<property>  
<name>mapreduce.framework.name</name>  
<value>yarn</value>  
</property>  
<property>  
<name>yarn.app.mapreduce.am.staging-dir</name>
```

```
<value>/user</value>
</property>
</configuration>
```

## 2. Replace your `/etc/hadoop/conf/yarn-site.xml` with below contents on master host

```
[root@master conf]# cat yarn-site.xml
```

```
<configuration>
<property>
<name>yarn.nodemanager.aux-services</name>
<value>mapreduce_shuffle</value>
</property>
<property>
<name>yarn.nodemanager.aux-services.mapreduce_shuffle.class</name>
<value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
<property>
<name>yarn.log-aggregation-enable</name>
<value>true</value>
</property>
<property>
<description>List of directories to store localized files in.</description>
<name>yarn.nodemanager.local-dirs</name>
<value>file:///var/lib/hadoop-yarn/cache/${user.name}/nm-local-dir</value>
</property>
<property>
<description>Where to store container logs.</description>
<name>yarn.nodemanager.log-dirs</name>
```

```
<value>file:///var/log/hadoop-yarn/containers</value>
</property>
<property>
<description>Where to aggregate logs to.</description>
<name>yarn.nodemanager.remote-app-log-dir</name>
<value>hdfs://master.hadoop.com:8020/var/log/hadoop-yarn/apps</value>
</property>
<property>
<description>Classpath for typical applications.</description>
<name>yarn.application.classpath</name>
<value>
$HADOOP_CONF_DIR,
$HADOOP_COMMON_HOME/*,$HADOOP_COMMON_HOME/lib/*,
$HADOOP_HDFS_HOME/*,$HADOOP_HDFS_HOME/lib/*,
$HADOOP_MAPRED_HOME/*,$HADOOP_MAPRED_HOME/lib/*,
$HADOOP_YARN_HOME/*,$HADOOP_YARN_HOME/lib/*
</value>
</property>
<property>
<name>yarn.resourcemanager.hostname</name>
<value>slave.hadoop.com</value>
</property>
<property>
<name>yarn.nodemanager.local-dirs</name>
<value>file:///data/1/yarn/local,file:///data/2/yarn/local,file:///data/3/yarn/local</value>
</property>
<property>
<name>yarn.nodemanager.log-dirs</name>
<value>file:///data/1/yarn/logs,file:///data/2/yarn/logs,file:///data/3/yarn/logs</value>
</property>
</configuration>
```

### 3. Copy modified files to slave machine.

```
[root@master conf]# scp mapred-site.xml yarn-site.xml slave.hadoop.com:/etc/hadoop/conf/
mapred-site.xml                                100%
1086   1.1KB/s  00:00
yarn-site.xml                                  100%
2787   2.7KB/s  00:00
[root@master conf]#
```

### 4. Configure local directories for yarn

To be done on yarn machine i.e. slave.hadoop.com in our case

```
[root@slave ~]# mkdir -p /data/1/yarn/local /data/2/yarn/local /data/3/yarn/local /data/4/yarn/local
[root@slave ~]# mkdir -p /data/1/yarn/logs /data/2/yarn/logs /data/3/yarn/logs /data/4/yarn/logs
[root@slave ~]# chown -R yarn:yarn /data/1/yarn/local /data/2/yarn/local /data/3/yarn/local /data/4/yarn/local
[root@slave ~]# chown -R yarn:yarn /data/1/yarn/logs /data/2/yarn/logs /data/3/yarn/logs /data/4/yarn/logs
```

## 5. Configure the history server.

Add below properties in mapred-site.xml

```
<property>
<name>mapreduce.jobhistory.address</name>
<value>master.hadoop.com:10020</value>
</property>
<property>
<name>mapreduce.jobhistory.webapp.address</name>
<value>master.hadoop.com:19888</value>
</property>
```

## 6. Configure proxy settings for history server

Add below properties in /etc/hadoop/conf/core-site.xml

```
<property>
<name>hadoop.proxyuser.mapred.groups</name>
<value>*</value>
</property>
<property>
<name>hadoop.proxyuser.mapred.hosts</name>
<value>*</value>
</property>
```



## 7. Copy modified files to slave.hadoop.com

```
[root@master conf]# scp mapred-site.xml core-site.xml slave.hadoop.com:/etc/hadoop/conf/
mapred-site.xml                                100
% 1299  1.3KB/s  00:00
core-site.xml                                  100% 1
174  1.2KB/s  00:00
[root@master conf]#
```

## 8. Create history directories and set permissions

```
[root@master conf]# sudo -u hdfs hadoop fs -mkdir -p /user/history
[root@master conf]# sudo -u hdfs hadoop fs -chmod -R 1777 /user/history
[root@master conf]# sudo -u hdfs hadoop fs -chown mapred:hadoop /user/history
```

## 9. Create log directories and set permissions

```
[root@master conf]# sudo -u hdfs hadoop fs -mkdir -p /var/log/hadoop-yarn
[root@master conf]# sudo -u hdfs hadoop fs -chown yarn:mapred /var/log/hadoop-yarn
```

## 10. Verify hdfs file structure

```
[root@master conf]# sudo -u hdfs hadoop fs -ls -R /
```

```
drwxrwxrwt - hdfs hadoop      0 2015-04-25 01:16 /tmp
drwxr-xr-x - hdfs hadoop      0 2015-04-25 02:52 /user
drwxrwxrwt - mapred hadoop     0 2015-04-25 02:52 /user/history
drwxr-xr-x - hdfs hadoop      0 2015-04-25 02:53 /var
drwxr-xr-x - hdfs hadoop      0 2015-04-25 02:53 /var/log
drwxr-xr-x - yarn mapred      0 2015-04-25 02:53 /var/log/hadoop-yarn
[root@master conf]#
```

## 11. Start yarn and Jobhistory server

### On slave.hadoop.com

```
[root@slave ~]# sudo service hadoop-yarn-resourcemanager start
starting resourcemanager, logging to /var/log/hadoop-yarn/yarn-yarn-resourcemanager-slave.hado
op.com.out
Started Hadoop resourcemanager:                [ OK ]
[root@slave ~]#
```

```
[root@slave ~]# sudo service hadoop-yarn-nodemanager start
starting nodemanager, logging to /var/log/hadoop-yarn/yarn-yarn-nodemanager-slave.hadoop.com.
out
Started Hadoop nodemanager:                    [ OK ]
[root@slave ~]#
```

### On master.hadoop.com

```

[root@master conf]# sudo service hadoop-mapreduce-historyserver start

starting historyserver, logging to /var/log/hadoop-mapreduce/mapred-mapred-historyserver-master
.hadoop.com.out

15/04/25 02:56:01 INFO hs.JobHistoryServer: STARTUP_MSG:

/*****

STARTUP_MSG: Starting JobHistoryServer

STARTUP_MSG: host = master.hadoop.com/192.168.111.130

STARTUP_MSG: args = []

STARTUP_MSG: version = 2.6.0-cdh5.4.0

STARTUP_MSG: classpath =

STARTUP_MSG: build = http://github.com/cloudera/hadoop -r c788a14a5de9ecd968d1e2666e8765
c5f018c271; compiled by 'jenkins' on 2015-04-21T19:18Z

STARTUP_MSG: java = 1.7.0_79

-
-
-

*****/

Started Hadoop historyserver:                [ OK ]

[root@master conf]#

```

## 12. Create user for running mapreduce jobs

```

[root@master conf]# sudo -u hdfs hadoop fs -mkdir /user/gaurav

[root@master conf]# sudo -u hdfs hadoop fs -chown kuldeep /user/gaurav

```

**13. Important: Don't forget to set core hadoop services to auto start when OS boot ups.**

**On master.hadoop.com**

```
[root@master conf]# sudo chkconfig hadoop-hdfs-namenode on  
[root@master conf]# sudo chkconfig hadoop-mapreduce-historyserver on
```

**On slave.hadoop.com**

```
[root@slave ~]# sudo chkconfig hadoop-yarn-resourcemanager on  
[root@slave ~]# sudo chkconfig hadoop-hdfs-secondarynamenode on  
[root@slave ~]# sudo chkconfig hadoop-yarn-nodemanager on  
[root@slave ~]# sudo chkconfig hadoop-hdfs-datanode on
```

**Final step : check UIs**

